

Comparing Hyperparameter Optimization Methods

Random Forest on the Titanic Dataset (Kaggle)

Machine Learning


October 5, 2025




Objective: Implement and compare three HPO methods for tuning a **Random Forest** on the Titanic dataset.

- Methods: **Hyperband, Bayesian Optimization, Genetic Algorithm.**
- Metrics reported on holdout: **Accuracy, F1, AUC.**
- Efficiency: **Search time** and **number of evaluations.**
- Robustness: **Stability** via $\text{std}(\text{F1})$ in cross-validation.


Dataset and Key Signals (EDA)



Class imbalance in the target.



Strong survival gap by Sex.



Pclass is highly predictive.

Experimental Protocol and Pipeline

- Split: **80/20 stratified** train/holdout; inner CV: **5-fold**.
- Preprocessing:
 - Numeric: `SimpleImputer(median)`.
 - Categorical: `SimpleImputer(most_frequent) + OneHotEncoder(handle_unknown=ignore)`.
- Model: `RandomForestClassifier(random_state=42, n_jobs=-1)`.
- Tuned hyperparameters: `n_estimators`, `max_depth`, `min_samples_split`, `max_features`.

Hyperband: Idea (Resource Allocation & Early Stopping)

- Evaluate many configurations with a small resource budget r (e.g., trees).
- Successive halving: keep the top fraction and *increase* r , pruning poor configs early.
- Multiple brackets trade off breadth vs. depth of search.
- Implementation: `HalvingRandomSearchCV` with resource `rf_n_estimators`.

Strengths: Very efficient when many bad configs exist; strong anytime performance.

Weaknesses: Random sampling inside each bracket; may miss narrow optima.

Bayesian Optimization: Idea (Surrogate & Acquisition)

- Fit a surrogate model $\hat{f}(\theta)$ of performance over hyperparameters θ .
- Use an acquisition function (e.g., Expected Improvement, UCB) to pick the next θ .
- Balances exploration (uncertain regions) and exploitation (promising regions).
- Implementation: BayesSearchCV (skopt) or TPE (Optuna).

Strengths: Good solutions with fewer evaluations.

Weaknesses: Surrogate assumptions; tuning of acquisition; overhead vs. random search.

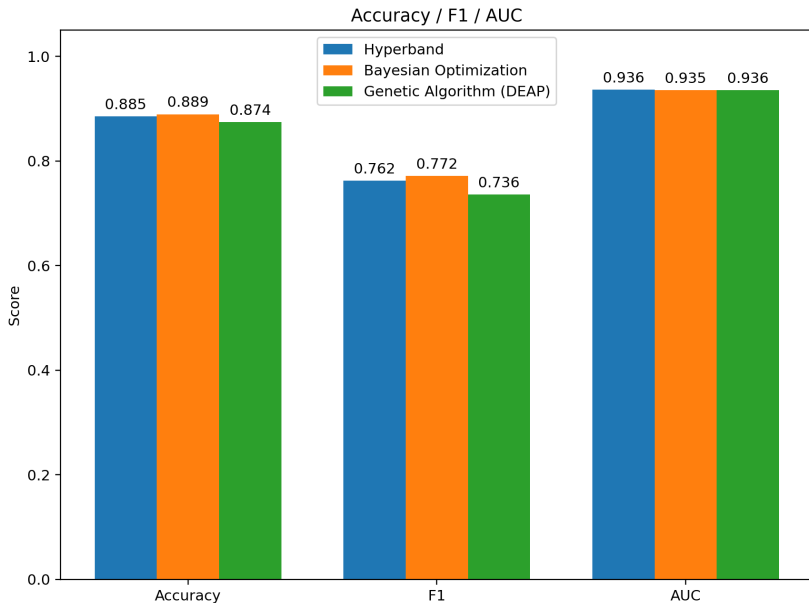
Genetic Algorithm: Idea (Population-Based Search)

- Population of candidate hyperparameters; selection by fitness (CV F1).
- Variation via **crossover** and **mutation**; iterate for G generations.
- Implementation: DEAP + `cross_val_score` with F1.

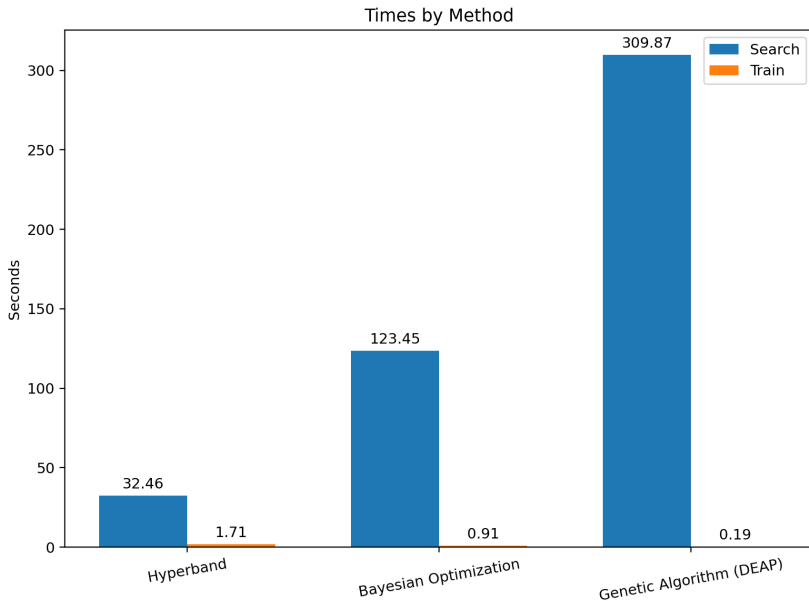
Strengths: Flexible, non-differentiable spaces; escapes local optima.

Weaknesses: Higher computational cost; multiple hyperparameters of the GA itself.

Holdout Performance



Efficiency: Time and Evaluations




Comparison Table (Summary)

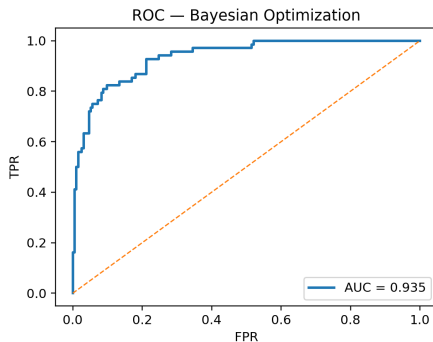
Method	Accuracy	F1	AUC	#Eval / Time (s)
Hyperband	0.885	0.762	0.936	22 / 32.46
Bayesian Optimization	0.889	0.772	0.935	32 / 123.45
Genetic Algorithm (DEAP)	0.874	0.736	0.936	215 / 309.87

Stability (std F1 in CV): *High* for Hyperband and Bayes; *Medium* for GA.

Best Method Diagnostics (Bayesian Optimization)



Confusion matrix: TN=184, FP=10,
FN=19, TP=49.



ROC curve on holdout; AUC = 0.935.

Which Metrics Matter Here?

- **F1**: balances precision/recall on the positive class (Survived) under class imbalance.
- **AUC**: global separability; similar (≈ 0.935) \Rightarrow focus on F1 to differentiate.
- **Computational cost**: search time and #evaluations for practical feasibility.
- **Stability**: low std(F1) indicates robustness across folds.

Conclusions

- **Bayesian Optimization** offers the best trade-off here: **highest F1 (0.772)** with moderate cost.
- **Hyperband** is preferred when **time budget is tight**: fast search, competitive F1.
- **Genetic Algorithm** explored widely but did **not** surpass F1 while being **computationally heavy**.